

# Computations with the IJC High-performance computing cluster

Angelika Merkel (Head of Bioinformatics)  
11/06/2024

# Connecting to the cluster

You can connect to the cluster (master node = minastirith) in several ways:

1. From anywhere with a browser via VPN Portal: <https://vpn.carrerasresearch.org> (bookmark 'Minastirith')
2. From your machine (Linux/Mac) with a terminal:  
\$ ssh username@minastirith
3. From your machine using a ssh client (e.g. [PUTTy for windows](#) ) connect directly to the host:  
minastirith

To connect from outside the IJC network you need to [connect to the SSHlogin node](#) first.

# Enter minastirth

\$ ssh amerkel@minastirth

```
[amerkel@IJC20249 ~]$ ssh minastirth
Last login: Mon Jun 10 17:06:04 2024 from 84.88.90.193

      o8o      -      o8o      o8o      .      oooo
      .o8      .o8      .o8      .o8      888
ooo. .oo. .oo. oooo ooo. .oo. .oooo. .oooo.o 888 oooo oood8b oooo 888 888
888P*Y88bP*Y88b`888`888P*Y88b`P`88b d88( "8 .o888oo`888 oooo d8b`888 .o888oo 888 .oo.
888 888 888 888 888 888 .oP`888`"Y88b. 888 888 888 888 888 888 .888 888
888 888 888 888 888 888 d8( 888 o. )88b 888 . 888 888 888 888 . 888 888
o888o o888o o888o o888o o888o o888o Y888 -8o 8 888P 888 o888o d888b o888o 888 o888o o888o

Master AMD EPYC 7272 12 cores 128Gb Ram *****
Node01 AMD EPYC 7702 128 cores 1Tb Ram *****
Node02 AMD EPYC 7702 128 cores 1Tb Ram *****
Node03 AMD EPYC 7702 128 cores 1Tb Ram Nvidia A40 GPU

Please remember:

1.- Save your stuff before it gets deleted! <3

      BEEGFS data (/mnt/beeufs) will be periodically removed.
      Transfer your data to /ijc/LABS or /ijc/USERS !

      Example: $ rsync -rlv --checksum source_folder destination_folder

2.- Do not launch jobs on the master node, use sbatch files or salloc. Interactive jobs on minastirth will be killed!

3.- Check performance of your finished jobs!

      Example: $ seff [jobid]

Thank you all!

Partition Name  Max Time  Default Time  Max CPU x Node  Max Mem x CPU  Def Mem x CPU  Max Mem x Node
=====
fast            2 hours   1h            64              4000M          2000M          256000M
highMem         1 Day     1h            8               -              -              800000M
hpc             3 Days    4h            80              200000M        4000M          -

Max run Cpu Per User: 48

show personal statistics
=====
module load reportseff
reportseff --since d=2 --format "+user,jobname,start,ncpus,reqmem" -u amerkel
[amerkel@minastirth ~]$
```

# Working with the HPC

input

software

instructions

resources

output


# Using the software

## Available as modules

Description	Command
see all available modules	<code>module avail</code>
load/unload a module	<code>module load/unload</code>
see all loaded modules	<code>module list</code>
unload all modules	<code>module purge</code>

# Executing a job

All tasks (jobs) executed on the cluster computational nodes are managed by SLURM (the scheduler)



<https://docs-research-it.berkeley.edu/services/high-performance-computing/user-guide/>

SLURM schedules each job based on available resources (CPU, memory, nodes, execution time, etc)

# SLURM script

`./my_clusterjob.sh # basic serial job`

```
#!/bin/bash

# SLURM arguments
#SBATCH --job-name=job_serial      # Job name
#SBATCH --cpus-per-task=1          # Run on a single CPU
#SBATCH --mem=4gb                  # Job memory request
#SBATCH --time=00:10:00            # Time limit hrs:min:sec
#SBATCH --output=job_%j.log        # Standard output and error log

# load software
module load R

# message something
echo "Running R script on a single CPU core"

# run
Rscript myscript.r
```

`%j = jobid`

# Slurm Commands

Command	Description
<code>sbatch my_clusterjob.sh</code>	Submit job for execution
<code>squeue</code>	Show the actual job queue
<code>scancel jobid</code>	Cancel a job
<code>sacct -j jobid</code>	Job accounting infos
<code>seff -j jobid</code>	Report on the efficiency of a job's cpu and memory utilization (after the job has finished)
<code>salloc</code>	Allocate resource for an interactive shell

`module load reportseff`

`reportseff --since d=2 --format "+user,jobname,start,ncpus,reqmem" -u username`



# Deploy interactive shell


You can deploy an interactive shell (with specified resource) on a computational node for testing:

```
$ salloc -c 4
```


```
[amerkel@minastirith ~]$ salloc
salloc: Granted job allocation 3216713
salloc: Waiting for resource configuration
salloc: Nodes c02 are ready for job
[amerkel@c02 ~]$
```

# Parallelization increases speed

1. Single input + single task



2. Multiple input + Single task



# Job arrays

```
#!/bin/bash

# SLURM arguments
#SBATCH --job-name=job_array      # Job name
#SBATCH --ntasks=1               # Run a single task
#SBATCH --mem=4gb                 # Job Memory
#SBATCH --time=00:05:00          # Time limit hrs:min:sec
#SBATCH --output=array_%A-%a.log # Standard output and error log
#SBATCH --array=1-5              # Array range

# message something
echo "This is task" $SLURM_ARRAY_TASK_ID

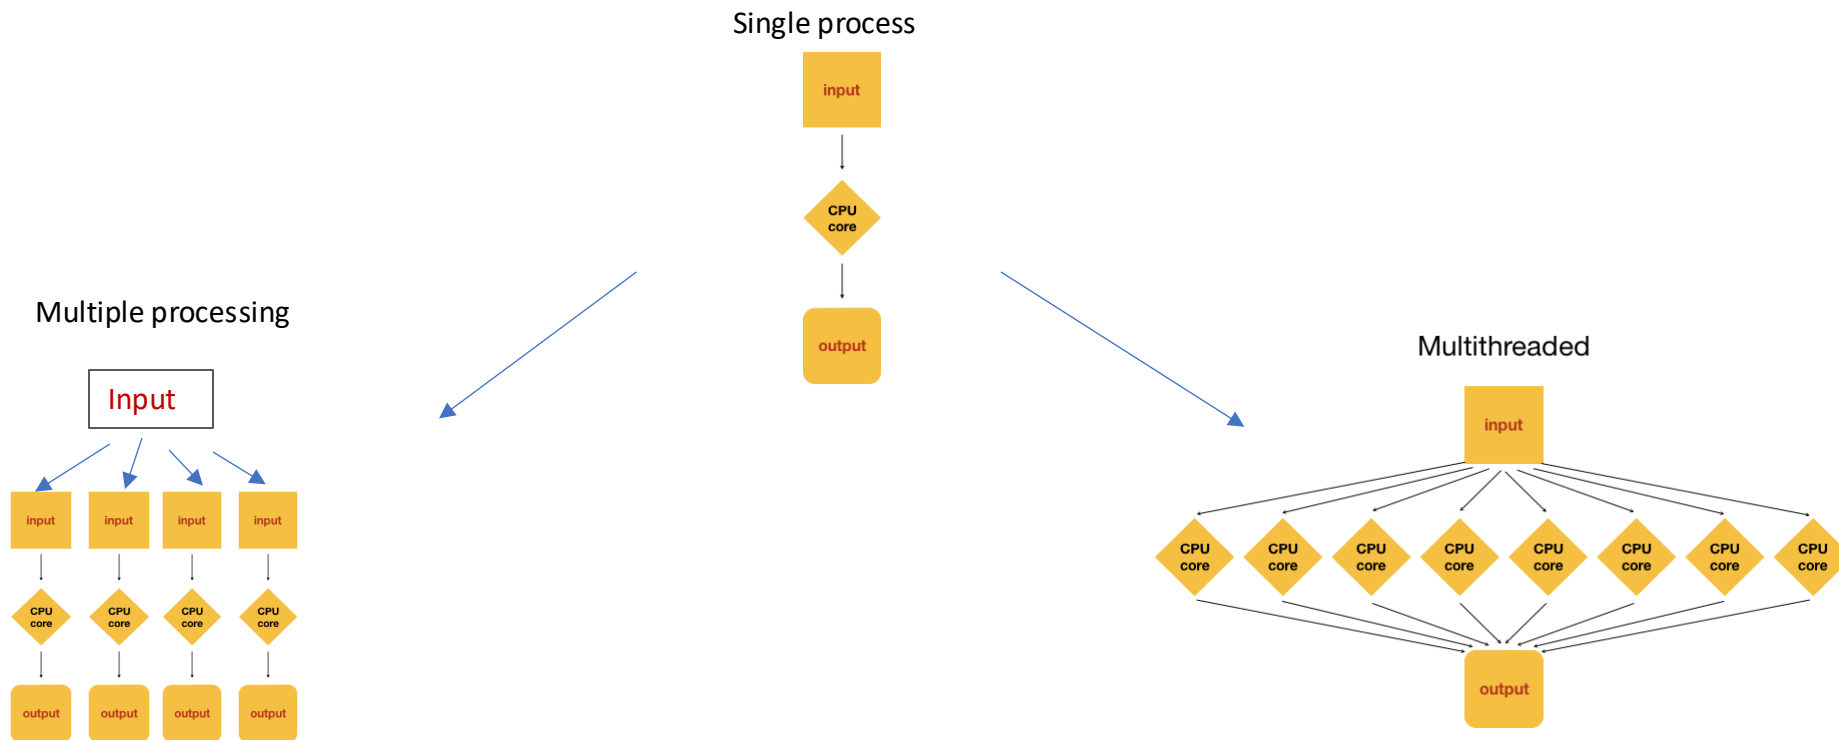
# run
infile=$(ls *.txt | sed -n ${SLURM_ARRAY_TASK_ID}p) # multiple input files (*.txt)

./myscript.sh $infile                                # run myscript for each input file
```

Job arrays allow to manage and submit similar jobs or execute the same job over multiple inputs (samples/data sets/parameter sets)


- Multiple jobs are created/submitted from one script, each with a different `$SLURM_ARRAY_TASK_ID` (=build-in variable)
- `$SLURM_ARRAY_TASK_ID` takes value specified with the array range

# Parallelization increases speed



# Multi-processing

Multiple processes




In **multi-processing** multiple processes are executed on multiple CPUs:

- `#SBATCH --ntask=number_of_processes`

# Multi-threading

Multithreaded



In **multi-threading** a single process is executed with multiple threads (e.g. multiple code segments):

1. Shared-memory multi-threading application (threaded, OpenMP, PTHREADS) can use multiple cpus but only on a single node:

- #SBATCH --nodes=1
- #SBATCH --ntask=1
- #SBATCH --cpus-per-task=number\_of\_threads

2. OpenMPI applications allow to share threads across nodes:

- #SBATCH --ntask=1
- #SBATCH --cpus-per-task=number\_of\_threads

# SLURM partitions (queues)

There are multiple partitions on the cluster depending to which jobs can be submitted depending on the required resources

#SBATCH --partition

Partition Name	Max Time	Default Time	Max CPU x Node	Max Mem x CPU	Def Mem x CPU	Max Mem x Node
fast	2 hours	1h	64	4000M	2000M	256000M
highMem	1 Day	1h	8	-	-	800000M
hpc	3 Days	4h	86	200000M	4000M	-

**\*NOTE:** If the resource requirement is not specified, default value = max value

# Do's & Don't's

## DO

- only request resources that you need (more CPUs or memories don't make your job faster!)
- check if the resources you requested were efficiently used (use `seff -j jobid` or similar)
- check if your job should be submitted to a special queue
- submit resource intensive jobs at low peak times (over night or the weekend)
- break large jobs down into smaller ones
- optimize your code

## DON'T

- ask for unnecessary resources (+ 10/15% time or memory are usually sufficient)
- execute any resource intensive task on the master node, use an interactive shell instead
- forget to close your interactive session



# How to get help

- contact **IT** for help with resource requirements, software and permissions
- contact **BIT** for help with coding and workflows
- check the **documentation** @:  
<https://wordpress.carrerasresearch.org/>  
<https://vpn.carrerasresearch.org/> > wordpress
- check the BIT workshop pages @:  
[https://ijcbit.github.io/Workshops/Trainings/HPC/Intro to IJC HPC.html](https://ijcbit.github.io/Workshops/Trainings/HPC/Intro_to_IJC_HPC.html)
- Teams: ijcbioinfo

Questions?

